

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv('mushrooms.csv')
df.head()
```

Out[2]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	stalk-surface-above-ring
0	p	x	s	n	t	p	f	c	n	k	...	s	s
1	e	x	s	y	t	a	f	c	b	k	...	s	s
2	e	b	s	w	t	l	f	c	b	n	...	s	s
3	p	x	y	w	t	p	f	c	n	n	...	s	s
4	e	x	s	g	f	n	f	w	b	k	...	s	s

5 rows × 23 columns



```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8124 entries, 0 to 8123
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   class                                8124 non-null   object
1   cap-shape                            8124 non-null   object
2   cap-surface                          8124 non-null   object
3   cap-color                            8124 non-null   object
4   bruises                              8124 non-null   object
5   odor                                 8124 non-null   object
6   gill-attachment                      8124 non-null   object
7   gill-spacing                         8124 non-null   object
8   gill-size                            8124 non-null   object
9   gill-color                           8124 non-null   object
10  stalk-shape                          8124 non-null   object
11  stalk-root                           8124 non-null   object
12  stalk-surface-above-ring             8124 non-null   object
13  stalk-surface-below-ring             8124 non-null   object
14  stalk-color-above-ring               8124 non-null   object
15  stalk-color-below-ring               8124 non-null   object
16  veil-type                            8124 non-null   object
17  veil-color                           8124 non-null   object
18  ring-number                          8124 non-null   object
19  ring-type                            8124 non-null   object
20  spore-print-color                    8124 non-null   object
21  population                           8124 non-null   object
22  habitat                              8124 non-null   object
dtypes: object(23)
memory usage: 1.4+ MB
```

```
In [4]: df.index
```

Out[4]: RangeIndex(start=0, stop=8124, step=1)

In [5]: `df.columns`

Out[5]: Index(['class', 'cap-shape', 'cap-surface', 'cap-color', 'bruises', 'odor', 'gill-attachment', 'gill-spacing', 'gill-size', 'gill-color', 'stalk-shape', 'stalk-root', 'stalk-surface-above-ring', 'stalk-surface-below-ring', 'stalk-color-above-ring', 'stalk-color-below-ring', 'veil-type', 'veil-color', 'ring-number', 'ring-type', 'spore-print-color', 'population', 'habitat'], dtype='object')

In [6]: `df['class'].value_counts()`

Out[6]: e 4208
p 3916
Name: class, dtype: int64

In [7]: `df.describe().transpose()`

Out[7]:

	count	unique	top	freq
class	8124	2	e	4208
cap-shape	8124	6	x	3656
cap-surface	8124	4	y	3244
cap-color	8124	10	n	2284
bruises	8124	2	f	4748
odor	8124	9	n	3528
gill-attachment	8124	2	f	7914
gill-spacing	8124	2	c	6812
gill-size	8124	2	b	5612
gill-color	8124	12	b	1728
stalk-shape	8124	2	t	4608
stalk-root	8124	5	b	3776
stalk-surface-above-ring	8124	4	s	5176
stalk-surface-below-ring	8124	4	s	4936
stalk-color-above-ring	8124	9	w	4464
stalk-color-below-ring	8124	9	w	4384
veil-type	8124	1	p	8124
veil-color	8124	4	w	7924
ring-number	8124	3	o	7488
ring-type	8124	5	p	3968
spore-print-color	8124	9	w	2388
population	8124	6	v	4040
habitat	8124	7	d	3148

```
In [8]: x=df.drop('class',axis=1)
x=pd.get_dummies(x,drop_first=True)
y=df['class']
```

```
In [9]: x.head()
```

```
Out[9]:
```

	cap- shape_c	cap- shape_f	cap- shape_k	cap- shape_s	cap- shape_x	cap- surface_g	cap- surface_s	cap- surface_y	cap- color_c	cap- color_e
0	0	0	0	0	1	0	1	0	0	0
1	0	0	0	0	1	0	1	0	0	0
2	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	1	0	0	1	0	0
4	0	0	0	0	1	0	1	0	0	0

5 rows × 95 columns

```
In [10]: y.head()
```

```
Out[10]: 0    p
1    e
2    e
3    p
4    e
Name: class, dtype: object
```

```
In [11]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)

model=LogisticRegression()
model.fit(x_train,y_train)

train_pred=model.predict(x_train)
test_pred=model.predict(x_test)
model.score(x_test,y_test)
```

```
Out[11]: 1.0
```

```
In [12]: test_pred=model.predict(x_test)
```

```
In [13]: model.score(x_test,y_test)
```

```
Out[13]: 1.0
```

```
In [14]: model.score(x_train,y_train)
```

```
Out[14]: 1.0
```

```
In [18]: from sklearn.model_selection import cross_val_score  
scores=cross_val_score(model,x,y,cv=5)  
scores.mean()
```

```
Out[18]: 0.8942316786661614
```

```
In [ ]:
```